

AeroAspire - SDE Intern

Gokul Krishna S

Week 4 – Day 3 (October 15)

Questions/Reflections:

1. How do you design efficient query for filtering? What is index; when do you use it?

- A filter in SQL narrows down the rows returned from a table. To make filtering efficient, you want to:
- Apply filters early — fetch only what you need. Let the WHERE clause or subquery do the trimming before any sorting or grouping.
- Avoid functions on filtered columns. For example, writing WHERE YEAR(hire_date) = 2024 prevents the database from using an index effectively. Instead, use range-based filtering:

- Sql :

```
SELECT * FROM employees
WHERE hire_date >= '2024-01-01' AND hire_date < '2025-01-01';
```

- Filter before sorting: if you sort a smaller dataset, you compute less. This can drastically improve query performance.
- Use equality (=) or range comparisons (>, <) instead of LIKE whenever possible. Comparing indexed columns with %LIKE% forces full scans.
- **What is an index and when do you use it?**
- An index is like a look-up table — a sorted data structure (like a book's index) that lets the database find specific rows faster without scanning the entire table.
- For example, an index on the email column means when you run:

- sql
SELECT * FROM users WHERE email = 'alex@example.com';
- the database jumps straight to matching rows instead of reading every user.
- Use an index when:
 - The column is frequently used in WHERE, JOIN, or ORDER BY clauses.
 - It contains highly selective values (many unique entries).
 - Use with care: indexes speed up reads, but slow down writes (INSERT, UPDATE, DELETE) since the index must update too. Also, indexes consume extra storage.

2. How does pagination work (offset/limit etc.).

- Pagination is how you fetch results in smaller chunks — common for APIs showing tables or lists.
- **Basic syntax:**
SELECT * FROM books ORDER BY published_at DESC
LIMIT 10 OFFSET 20;
- LIMIT: number of rows to return.
- OFFSET: number of rows to skip first.
- **For example:**
Page 1 → LIMIT 10 OFFSET 0
Page 2 → LIMIT 10 OFFSET 10
- Performance tip: large offsets are slow because the database still scans skipped rows. An alternative is cursor or keyset pagination:
- sql
SELECT * FROM books
WHERE id > 1000 ORDER BY id ASC LIMIT 10;
- This uses an indexed column (id) for pagination and avoids skipping.

3. What's the flow of building these endpoints, receiving query params, applying them in SQL / ORM, returning results.

- When you create an endpoint like `/api/products`, the backend usually:
- **Receives query parameters from the client:**
`/api/products?category=shirts&limit=10&page=2.`
- **Parses and validates them in your backend code:**

```
const { category, limit = 10, page = 1 } = req.query;  
const offset = (page - 1) * limit;
```
- **Builds a parameterized SQL or ORM query:**

```
SELECT * FROM products  
WHERE category = $1  
ORDER BY created_at DESC  
LIMIT $2 OFFSET $3;
```
- **or with an ORM like Sequelize:**

```
Product.findAll({  
  where: { category },  
  order: [['createdAt', 'DESC']],  
  limit,  
  offset  
});
```
- Executes the query — the database handles filtering, sorting, and pagination efficiently when columns are properly indexed.
- **Returns a structured response:**

```
{  
  "page": 2,  
  "limit": 10,  
  "total": 125,  
  "data": [ ...products ]  
}
```